

# **Computational Data Analysis**

## **The Support Vector Machine and Convex Optimization**

Lars Arvastson and Line Clemmensen

February 28, 2018

# Today's Lecture

- ▶ Recap
- ▶ Convex optimization using Lagrange multipliers
- ▶ Optimal Separating Hyperplanes
- ▶ Support Vector Machines
- ▶ The kernel trick

# Last Week

- ▶ Linear discriminant analysis and Logistic regression
  - ▶ What for?
  - ▶ How do they compare?
- ▶ Basis expansion
  - ▶ What is it?
  - ▶ How did we use it?



# Crash course in constrained optimization

We learn to solve  $\begin{cases} \max_x f(x) \\ g(x) = 0 \\ h(x) \geq 0 \end{cases}$

using Lagrange multipliers

## Why?

Because we will use it to build the Support Vector Machine!

# Unconstrained optimization

Solve

$$\max_x f(x)$$

Assume that  $f$  is nice, i.e. continuously differentiable.

Then a local maxima,  $x^*$  fulfills

1. Gradient is zero,  $\nabla_x f(x^*) = 0$
2. Hessian is negative definite,  $v^T \nabla_{xx}^2 f(x^*) v < 0, \forall v \in \mathbb{R}^n$

where


$$\nabla_x f(x) = \begin{pmatrix} \frac{\partial f(x)}{\partial x_1} \\ \vdots \\ \frac{\partial f(x)}{\partial x_n} \end{pmatrix} \quad \nabla_{xx}^2 f(x) = \begin{pmatrix} \frac{\partial^2 f(x)}{\partial x_1^2} & \cdots & \frac{\partial^2 f(x)}{\partial x_1 \partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f(x)}{\partial x_n \partial x_1} & \cdots & \frac{\partial^2 f(x)}{\partial x_n^2} \end{pmatrix}$$

A negative second derivative guarantees a **local maximum** (otherwise saddle point or local minimum).

# Constrained optimization

Now, assume that any  $x$  is not good enough. Introduce a constraint that  $x$  must fulfill,

$$\begin{cases} \max_x f(x) \\ g(x) = 0 \end{cases}$$



- The stationary points are defined by  $\nabla f = -\lambda \nabla g$  for some constant  $\lambda$

# Lagrange multipliers

Define the **Lagrange primal function**

$$L_p(x, \lambda) = f(x) + \lambda g(x)$$

and the **Lagrange multiplier**  $\lambda$ .


Find solution  $(x^*, \lambda^*)$  to

$$\max_x \min_{\lambda} L_p(x, \lambda)$$

by solving


$$\begin{cases} \frac{\partial L_p}{\partial x} = 0 \\ \frac{\partial L_p}{\partial \lambda} = 0 \end{cases} \quad \text{i.e. } \nabla L_p = 0.$$

The stationary points,  $x^*$ , might be local maxima, local minima or saddle points. Verify that the Hessian is negative semi-definite.



## Example

$$\begin{cases} \max_x f(x_1, x_2) = 1 - x_1^2 - x_2^2 \\ g(x_1, x_2) = x_1 + x_2 - 1 = 0 \end{cases}$$



$$\begin{aligned} L_P(\mathbf{x}, \lambda) &= f(x) + \lambda g(x) \\ &= 1 - x_1^2 - x_2^2 + \lambda(x_1 + x_2 - 1) \end{aligned}$$

$$\begin{cases} \frac{\partial L_P}{\partial x_1} = -2x_1 + \lambda = 0 \\ \frac{\partial L_P}{\partial x_2} = -2x_2 + \lambda = 0 \\ \frac{\partial L_P}{\partial \lambda} = x_1 + x_2 - 1 = 0 \end{cases}$$

Solution/optimum is at

$(x_1^*, x_2^*) = (1/2, 1/2)$  with  $\lambda = 1$ .




# Inequalities in constraints, $h(x) \geq 0$

Constrained optimization with inequality constraints

$$\begin{cases} \max_x f(x) \\ h(x) \geq 0 \end{cases}$$

The optimum is either within the feasible region  $h(x) \geq 0$  or along the edge  $h(x) = 0$ .



Notice that  $\nabla h$  is always pointing inwards since  $h > 0$  in the feasible region and  $h = 0$  along the edge.

# Inequalities in constraints, $h(x) \geq 0$

## 1. Constraint is inactive, ie $h(x_1) > 0$ .

- Solution by  $\nabla f(x_1) = 0$ , ie Lagrange function,

$$\begin{cases} \nabla f(x_1) = -\mu \nabla h(x_1) \\ \mu = 0 \end{cases}$$


- Notice that  $\mu h(x_1) = 0$ , since  $\mu = 0$
- Maximum if negative definite Hessian

## 2. Constraint is active, ie $h(x_2) = 0$ .

- As before with  $\mu \neq 0$ . Important with the sign of  $\mu$ . In maximum  $\nabla f(x_2)$  is pointing out of the region  $h(x_2) > 0$ , ie

$$\begin{cases} \nabla f(x_2) = -\mu \nabla h(x_2) \\ \mu > 0 \end{cases}$$

- Notice that  $\mu h(x_1) = 0$ , since  $h(x_2) = 0$
- Maximum if negative semidefinite Hessian



# Lagranges problem for inequality constraints

A local maximum to the constrained optimization problem

$$\begin{cases} \max_x f(x) \\ h(x) \geq 0 \end{cases}$$

with Lagrange function

$$L_p(x, \mu) = f(x) + \mu h(x)$$

is given by  $(x^*, \mu^*)$  when (Karush-Kuhn-Tucker conditions)

1.  $\nabla_x L_p(x^*, \mu^*) = 0$
2.  $\mu^* \geq 0$
3.  $\mu^* h(x^*) = 0$
4.  $h(x^*) \geq 0$
5. Negative definite constraints on Hessian

For a minimization problem we change sign,  $L_p(x, \mu) = f(x) - \mu h(x)$ .

# Multiple constraints

Multiple constraints,

$$\begin{cases} \max_x f(x) \\ g_j(x) = 0 & \forall j \\ h_k(x) \geq 0 & \forall k \end{cases}$$

are handle with more Lagrange multipliers,

$$L_p(x, \lambda, \mu) = f(x) + \sum_j \lambda_j g_j(x) + \sum_k \mu_k h_k(x).$$

# Lagrange dual problem

The **Lagrange primal** problem is

$$\max_x \min_{\substack{\lambda \\ \mu \geq 0}} L_P(x, \lambda, \mu)$$

If we swap the order of min and max we get the **Lagrange dual** problem,

$$\min_{\substack{\lambda \\ \mu \geq 0}} \max_x L_P(x, \lambda, \mu)$$

Often these two problems have the same solution.

Define **Lagrange dual function**

$$L_D(\lambda, \mu) = \max_x L_P(x, \lambda, \mu)$$


# Slater's condition

## Slater's condition


The primal and dual optimization problems are equivalent when  $f$  is concave and constraints are convex.

- ▶ There must be some  $x$  fulfilling all constraints
- ▶ Linear constraints are OK
- ▶ A local optimum will also be the global optimum.
- ▶ Not necessary to check conditions on the Hessian.


Convex function




Concave function




Convex set



Non-convex set



## Example - Shortest distance from point to line



$$\begin{cases} \arg \min_x \frac{1}{2}(x_i - x)(x_i - x)^T \\ \text{such that} \\ x\beta + \beta_0 = 0 \end{cases}$$

Solve using Lagrange primal function



# Optimal Separating Hyperplane

- ▶ Binary classification
- ▶ Sometimes data are perfectly separated by a straight line
- ▶ No overlap, one class on one side and the other class on the other side
- ▶ Not very useful in practice but it can be modified into the powerful Support Vector Machine



# The decision function

## Linear decision functions

$$y_{\text{new}} = \text{sign}(x_{\text{new}}\beta + \beta_0)$$


- ▶ For practical reason we label the two classes 1 and  $-1$ .
- ▶ Fitting the model involves choosing values for  $\beta$  and  $\beta_0$
- ▶ Binary classification, extensions can be made
  - ▶ One vs. the rest
  - ▶ One vs. one

both approaches uses several models.



# The decision function

- ▶ Many hyperplanes can separate the two classes
- ▶ What would be optimal?



## **Linear Discriminant Analysis,**

used all data to define  $\Sigma$  and  $\mu$  from which the decision line was derived.

## **Logistic regression,**

defined decision line emphasizing data close to line.

## **Optimal Separating Hyperplanes,**


goes to the extreme and defines decision line based on closest observations only.

# Introduce the margin

Maximize the distance  $C$  from the decision line to the nearest points in each class.

There is **no probabilistic model** here as we have for linear discriminant analysis and logistic regression


Hence, no probability for class belonging and no ML-estimation



# Distance from point to plane

We wish to maximize the margin between classes.

We need an expression for point-to-plane distance




# OSH as a maximization problem

We can now formulate a maximization problem

$$\arg \max_{\beta, \beta_0} C$$

such that

$$y_i \frac{x_i \beta + \beta_0}{\|\beta\|} \geq C \quad \forall i$$



# The margin

Let  $x_+$  be a support point in class 1 and  $x_-$  be a support point in class -1. Then

$$C = \frac{1}{2} \frac{\beta^T}{\|\beta\|} (x_+ - x_-)$$



Why?

Margin  $C$  is invariant to length of  $\beta$ . Choose length of  $\beta$  such that

$$\beta^T x_+ + \beta_0 = 1$$

$$\beta^T x_- + \beta_0 = -1$$

which gives

$$\beta^T (x_+ - x_-) = 2$$

The margin becomes

$$C = \frac{1}{\|\beta\|}$$

and the constraints simplifies into

$$y_i(x_i\beta - \beta_0) \geq 1 \quad \forall i$$

# Solving the OSH problem

- ▶ Maximization problem can be turned into a minimization problem

$$\left\{ \begin{array}{l} \arg \min_{\beta, \beta_0} \frac{1}{2} ||\beta||^2 \\ \text{such that} \\ y_i(x_i\beta + \beta_0) \geq 1 \quad \forall i \end{array} \right.$$

- ▶ This is a nonlinear problem with **linear** constraints
  - ▶ You could use Matlabs `fmincon` function for constrained optimization of **any** nonlinear function
  - ▶ But this one is quadratic (convex) - does this simplify things?
- ▶ Quadratic programming
  - ▶ Very efficient solvers exists
  - ▶ Matlabs `quadprog`

# Dual formulation of OSH problem

We have formulated the OSH problem such that we can use efficient standard numerical solvers. **We have**

- ▶ A model with one  $\beta$  coefficient for each dimension of  $x$ .
- ▶ One constraint for each observation  $x$
- ▶ An optimal linear separation between classes

**What else** could we ask for? Well, it would be nice with


- ▶ One coefficient for each observation instead of each dimension.
  - ▶ Good idea for high-dimensional problems with few observations.
- ▶ A non-linear separation between classes.

We can achieve this if we use **Lagrange multipliers**



# Solving the OSH problem

Use the Lagrange multipliers!



$$\left\{ \begin{array}{l} \arg \min_{\beta, \beta_0} \frac{1}{2} \|\beta\|^2 \\ \text{such that} \\ y_i(x_i\beta + \beta_0) \geq 1 \quad \forall i \end{array} \right.$$

**Step 1** Incorporate constraints using Lagrange multipliers,  $\alpha_i$ ,

$$\left\{ \begin{array}{l} L(\beta, \beta_0, \alpha) = \frac{1}{2} \|\beta\|^2 - \sum_{i=1}^n \alpha_i (y_i(x_i\beta + \beta_0) - 1) \\ \alpha_i \geq 0 \quad \forall i \end{array} \right.$$

# Solving the OSH problem, cont'd

**Step 2** Differentiate and set to zero. This solves  $\arg \min_{\beta, \beta_0} L_p$  (Lagrange dual),

$$\begin{cases} \frac{\partial L}{\partial \beta} = \beta - \sum_i \alpha_i y_i x_i^T = 0 \\ \frac{\partial L}{\partial \beta_0} = \sum \alpha_i y_i = 0 \end{cases}$$

and we have

$$\begin{cases} \beta = \sum_i \alpha_i y_i x_i^T \\ \sum \alpha_i y_i = 0 \end{cases}$$

# OSH dual formulation

**Step 3** Plug into original problem and simplify

$$\begin{aligned}L_D &= \frac{1}{2} \left\| \sum \alpha_i y_i \mathbf{x}_i^T \right\|^2 - \sum (\alpha_i y_i (\mathbf{x}_i^T \beta + \beta_0) - \alpha_i) \\&= \dots \\&= \sum \alpha_i - \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \mathbf{x}_i \mathbf{x}_j^T \\&= \alpha^T \mathbf{1} - \frac{1}{2} \alpha^T Y X X^T Y \alpha \quad \text{where } Y = \text{diag}(y)\end{aligned}$$

This is Lagrange dual function. Dual formulation is OK since quadratic function with linear constraints fulfills **Slater's** conditions.

# OSH dual formulation, cont'd

**Step 4** Identify the QP components

$$\left\{ \begin{array}{l} \arg \max_{\alpha} \alpha^T \mathbf{1} - \frac{1}{2} \alpha^T Y X X^T Y \alpha \\ \text{such that} \\ \alpha_i \geq 0 \quad \forall i \\ \sum \alpha_i y_i = 0 \end{array} \right.$$

The general form of a QP problem is

$$\left\{ \begin{array}{l} \arg \min_{\alpha} \alpha^T Q \alpha + c^T \alpha \\ \text{such that} \\ A \alpha \leq b \\ E \alpha = d \end{array} \right.$$

- Identify  $Q$ ,  $c$ ,  $A$ ,  $b$ ,  $E$  and  $d$ ?
- How do we get  $\beta$ ?



# Two more things...

## How do we find $\beta_0$ ?


- ▶ For the support points we have  $y_i(x_i\beta + \beta_0) = 1$
- ▶ Use one of the support points to calculate  $\beta_0$

## How do we predict class belonging?

- ▶ Along the decision line we have  $x\beta + \beta_0 = 0$
- ▶ Along the support lines we have  $x\beta + \beta_0 = \pm 1$
- ▶ Decision based on the side of the line  
 $\hat{y}(x_{\text{new}}) = \text{sign}(x_{\text{new}}\beta + \beta_0)$

# Wait a minute... Margin??

- ▶ For overlapping data, there is no solution
  - ▶ What's the use?



- ▶ Can be modified into the **Support Vector Machine**
  - ▶ Handles overlapping observations.
  - ▶ Kernel trick for non-linear data.

# Support Vector Machine

- ▶ Most classification problems have overlapping classes.
- ▶ Let us modify the OSH such that we allow for some overlap
- ▶ This is the **Support Vector Machine**
- ▶ Used together with the kernel trick SVM is one of our most flexible classifiers


# SVM Cost Function

We got OSH from

$$\begin{cases} \arg \min_{\beta, \beta_0} \frac{1}{2} \|\beta\|^2 \\ \text{such that} \\ y_i(x_i\beta - \beta_0) \geq 1 \quad \forall i \end{cases}$$

Now, allow some overlap

$$\begin{cases} \arg \min_{\beta, \beta_0} \frac{1}{2} \|\beta\|^2 + \lambda \sum_{i=1}^n \xi_i \\ \text{such that} \\ y_i(x_i\beta - \beta_0) \geq 1 - \xi_i \quad \forall i \\ \xi_i \geq 0 \quad \forall i \end{cases}$$



We give our self a **budget for overlap**.

Smaller budget - larger  $\lambda$  - noisier solution



# Solving the SVM Problem

Similar to OSH

$$\left\{ \begin{array}{l} \arg \min_{\beta, \beta_0} \frac{1}{2} \|\beta\|^2 + \lambda \sum_{i=1}^n \xi_i \\ \text{such that} \\ y_i(x_i\beta + \beta_0) \geq 1 - \xi_i \quad \forall i \\ \xi_i \geq 0 \quad \forall i \end{array} \right.$$

Lagrange multiplier, differentiate, plug back...

$$\left\{ \begin{array}{l} \arg \max_{\alpha} \alpha^T \mathbf{1} - \frac{1}{2} \alpha^T Y X X^T Y \alpha \\ \text{such that} \\ \sum \alpha_i y_i = 0 \\ 0 \leq \alpha_i \leq \lambda \quad \forall i \end{array} \right.$$

# Comparison with OSH

**Optimal separating hyperplanes**


$$\left\{ \begin{array}{l} \arg \max_{\alpha} \alpha \mathbf{1} - \frac{1}{2} \alpha^T Y X X^T Y \alpha \\ \text{such that} \\ \alpha_i \geq 0 \quad \forall i \\ \sum \alpha_i y_i = 0 \end{array} \right.$$

**Support vector machine**

$$\left\{ \begin{array}{l} \arg \max_{\alpha} \alpha \mathbf{1} - \frac{1}{2} \alpha^T Y X X^T Y \alpha \\ \text{such that} \\ 0 \leq \alpha_i \leq \lambda \quad \forall i \\ \sum \alpha_i y_i = 0 \end{array} \right.$$

Both are quadratic programming problems with linear constraints

# Comparison with logistic regression



With  $f(x) = x\beta + \beta_0$  and  $y_i \in \{-1, 1\}$ , consider

$$\min_{\beta, \beta_0} \sum_{i=1}^N (1 - y_i f(x_i))_+ + \frac{\lambda}{2} \|\beta\|^2$$

This hinge loss criterion is equivalent to the SVM.

Compare with

$$\min_{\beta, \beta_0} \sum_{i=1}^N \log(1 + e^{-y_i f(x_i)}) + \frac{\lambda}{2} \|\beta\|^2$$

(In Lecture 3 we used  $y_i \in \{0, 1\}$ .)

This is the ML formulation of **ridged logistic regression**

# Basis expansion and kernels

- ▶ We can do SVM (and OSH) on a transformed feature space
- ▶ Transformed features gives non-linear decision boundaries.
- ▶ With **the Kernel trick** we can use an infinite dimensional feature expansion

# Non-linear SVM

Let's try basis expansions!

$$\left\{ \begin{array}{l} \arg \max_{\alpha} \alpha \mathbf{1} - \frac{1}{2} \alpha^T Y X X^T Y \alpha \\ \text{such that} \\ 0 \leq \alpha_i \leq \lambda \quad \forall i \\ \sum \alpha_i y_i = 0 \end{array} \right.$$

Use  $h(X)$  instead of  $X$ ,

$$\left\{ \begin{array}{l} \arg \max_{\alpha} \alpha \mathbf{1} - \frac{1}{2} \alpha^T Y h(X) h(X)^T Y \alpha \\ \text{such that} \\ 0 \leq \alpha_i \leq \lambda \quad \forall i \\ \sum \alpha_i y_i = 0 \end{array} \right.$$

- ▶  $h(X) : R^p \rightarrow R^M$ , e.g.  $[x_1 \ x_2] \rightarrow [x_1 \ x_2^2 \ x_1 x_2]$
- ▶  $h(X)h(X)^T$  is of size  $n \times n$

# The kernel trick

The term  $h(X)h(X)^T$  does not depend on  $M$ , the number of basis functions.

We only need to specify  $K(X)$  such that  $h(X)h(X)^T = K(X)$  - we call  $K$  a **kernel**. Then  $h$  is implicitly defined by  $K$ .

Common kernels

**Polynomial**  $K_{i,j} = (1 + x_i x_j^T)^d$  ( $x_i$  is observation  $i$ , ie row  $i$  in  $X$ )

**Radial**  $K_{i,j} = \exp\left(-\frac{1}{c} \|x_i - x_j\|^2\right)$

**Gaussian**  $K_{i,j} = \exp\left(-\frac{1}{2\sigma^2} \|x_i - x_j\|^2\right)$

**Neural network**  $K_{i,j} = \tanh(c_1 x_i x_j^T + c_2)$

# SVM with kernels

The new optimization problem

$$\left\{ \begin{array}{l} \arg \max_{\alpha} \alpha^T \mathbf{1} - \frac{1}{2} \alpha^T Y K Y \alpha \\ \text{such that} \\ 0 \leq \alpha_i \leq \lambda \quad \forall i \\ \sum \alpha_i y_i = 0 \end{array} \right.$$

To classify a new observations

$$\begin{aligned} \hat{y}_{\text{new}} &= \text{sign}(\beta h(x_{\text{new}}) + \beta_0) \\ &= \text{sign} \left( \sum_{i=1}^n \alpha_i y_i K(x_{\text{new}}, x_i) + b_0 \right) \end{aligned}$$

Calculate  $b_0$  using one of the points,  $i$ , on the margin,

$$b_0 = y_i - \sum_{j=1}^n \alpha_j y_j K(x_i, x_j)$$

# Phew!

- ▶ We have found an efficient way of maximizing the margin between classes
- ▶ Of course, there are software packages for you!





# SVM in Matlab

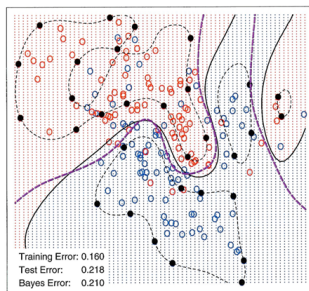

```
SVModel = fitcsvm(X,Y,Name,Value);  
Yhat = predict(SVModel,X);
```

Name	Value
'BoxConstraint'	$\lambda$ (Inf for OSH.)
'KernelFunction'	'linear', 'rbf', 'polynomial'
'KernelScale'	$c$
'PolynomialOrder'	positive integer
'Standardize'	good idea for SVM also

Use 'OptimizeHyperparameters' to select parameters for tuning and 'HyperparameterOptimizationOptions' to define a grid-search and cross validation.

# Example

Linear SVM and enlarged feature space using RBF kernel



# Model selection and SVM

## Use SVM together with Radial Basis Function kernel

- ▶  $K_{i,j} = \exp\left(-\frac{1}{c}||x_i - x_j||^2\right)$
- ▶ This give one parameter  $c$

## From the SVM loss function

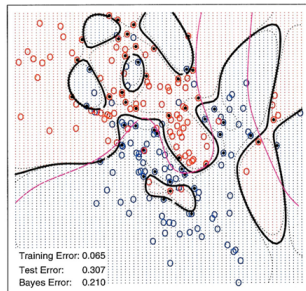

- ▶  $\arg \min_{\beta, \beta_0} \frac{1}{2} ||\beta||^2 + \lambda \sum_{i=1}^n \xi_i$
- ▶ This gives another parameter  $\lambda$

## Select parameters using cross validation

- ▶ Extensive search for  $c$
- ▶  $\lambda$  less crucial, try different values

# Overfitting

Overfitting is easy in an enlarged feature space!



# Caveats

- ▶ Kernel methods do not scale well. Limited to around 10000-20000 observations
- ▶ Kernel methods do not do variable selection in any reasonable or automatic way
  - ▶ With more features than observations there is always a separating hyperplane
  - ▶ Actually infinitely many which we have to choose between
- ▶ Potential problem with large number of features if many of them are garbage
- ▶ SVM do not generalize gracefully when the number of classes are more than two
  - ▶ Frequently used for multiclass classification anyway

# Summary

- ▶ Constrained optimization
  - ▶ Lagrange multipliers
  - ▶ Primal and dual formulation
- ▶ Optimal separating hyperplanes
  - ▶ Margin
  - ▶ Support vectors and support points
- ▶ Support vector machine
  - ▶ Budget for overlap
  - ▶ Kernel trick

**Questions?**